




Claude

Deep-dive

The Ultimate Workshop for Claude - 22/04/26

We explore the different tools offered by Anthropic and hopefully inspire you along the way to build something amazing

Ella Garcia | Rishi Kashyap



Today's Topics

1. The Claude App Ecosystem
 2. Settings Walkthrough & Chat+Artifacts
 3. Token Usage
 4. Common Security Issues
 5. Projects
 6. CoWork Deepdive
 - a. Connectors, Plugins and MCP
 - b. Dispatch
 - c. Scheduled Tasks
 7. Design
 8. Claude in Chrome
 9. Managed Agents and API
 10. Dashboards
 11. Claude Code; Plan Mode
 12. Claude for Office 365
 13. Claude Platform
-



ANTHROPIC

Founded 2021 by Dario & Daniela Amodei
(siblings) + 5 other OpenAI defectors

\$380B valuation. \$30B Annual
Revenue

\$0 → \$30B revenue in ~3 years = one of the
fastest revenue ramps in software history

Anthropic just passed OpenAI in annual revenue (April 7, 2026)
With just 5% of ChatGPT's user
base



The Rise of Claude Ecosystem

Chat

What you are used to.

Regular chat interface with
Artifacts

Limit on number of files and file
size

Context Window can get filled up

Cowork

Bypass all of the limitations from
Chat

Access to file system to easily
use documents

Work on complex tasks for
extended periods without
conversation timeouts or context
limits interrupting your progress.

Code

Coding Agent

Has many skills and tools to help
developers write code

Superior Handling of agents to
write, test and debug code

Can be used for non-coding
tasks like notes, brainstorming

Projects

Walkthrough of Settings + Chat

Claude App has many settings that are worth looking at before you start working on projects or tasks



Token Usage

Working with a limited plan - try to optimize token usage so that you don't hit your limits. Its always a **tradeoff with quality**

Choosing the right model

Haiku vs Sonnet vs Opus

Effective Prompting and output style

Do more with less talking. RICO.

Images

Images are tokenized by dimensions, not content - a 1080p screenshot can cost 1,500–4,000 tokens.

Files

Uploading a large PDF or document loads the entire content into context

Conversation length

Every turn re-sends the full history


System Prompts & Projects

They get loaded on every chat. Be crisp

Chain Inputs

"Amazing. Now do this". "Perfect, now lets do this".. XXX
INSTEAD

"Amazing, now do these 3 things. 1, 2 and 3 and check for abc"



Common Ways to Get Hacked

WATCH OUT - AI cannot save you every time. Yet.

Credential leakage: API keys, passwords, or tokens

Protect them at any cost

Prompt injection


AI can fall for prompts. Includes websites that are not safe

Sketchy Connectors/Plugins/MCP/Skills

Verify 3rd party tools that connect to your system. Non-trusted ones in particular

status.claude.com

<https://claude.com/blog>



Projects

A smart way to get tasks done. Skills make it powerful

Memory

Instructions

Files

Skills

Chats



Cowork Deepdive

Connectors/Plugins/MCP


Dispatch

Scheduled
Tasks

Ensure

- 1) Guardrails that tell Claude not take any dangerous actions without your permission
- 2) Access to folders that are necessary

<https://support.claude.com/en/articles/13345190-get-started-with-claude-cowork>



Claude Design

UI screens, prototypes , slides




<https://claude.ai/design>



Released April 17, 2026

Generate UI mockups, presentation slides, and prototypes directly from a text prompt.



Claude in Chrome



<https://fill.dev/form/identity-simple>



Managed Agents

Scheduled AI agent tasks that run on Anthropic's cloud infrastructure

Why it matters

Your laptop can be closed. Your phone in your pocket. The agent keeps running on Anthropic's servers and delivers results.

Real use cases

Every morning at 7am: summarize overnight Slack + email → send digest

Every hour: scrape competitor pricing → update Google Sheet

Access & limits (research preview)

Pro (\$20/mo): 5 runs/day

Max: 15 runs/day

Requires Claude Code on web enabled

<https://www.anthropic.com/engineering/managed-agents>

Claude API

**70% of Fortune 100 companies using it as of
2026**

Claude API is what drove Anthropic's rise in revenue as Coding Agents integrated Claude Models via API

<https://platform.claude.com/docs/en/api/overview>

Endpoint -

<https://api.anthropic.com>



If we have time:

- Dashboards
- Claude Code + Plan Mode
- Claude for Office 365
- Claude Platform





Presented by: Ella, Rishi & Sarfaraz

OPENCLAW101 + DEMO PART 1

Exploring the magic of 24/7 agents

OpenClaw changed the way we think of interacting with AI system

♥ 300k+





TODAY'S TOPICS

- Chat vs Agents
- Intro to OpenClaw
- Understanding Tokens & APIs
- Security
- System Design
- Baseline / Core Foundation
- Costs
- Overview of next tracks
- Setup (VM/VPS)

CHAT VS AGENTS




Chat

Human-in-the-loop, single-turn or multi-turn conversation. Context is determined solely by the user

User drives the interaction turn by turn. The model responds but doesn't do anything outside the conversation

Gemini, Custom GPTs etc.



Agent

Goal-oriented, autonomous, ^S can use tools and take actions with minimal human intervention

The model is given a goal and can plan, use tools (APIs, browsers, code execution), loop on its own, and take real-world actions.

n8n, Claude CoWork, Perplexity Computer etc.

2022

2024

2025

2026

ChatGPT Launch

Chat based applications

Agentic AI

Orchestration
of Agents

WHAT IS OPENCLAW?



Created by Peter Steinberger first in Nov 2025, was called Moltbot → Clawdbot → now OpenClaw



Runs on Your Machine

Mac, Windows, or Linux. Anthropic, OpenAI, or local models. Private by default—your data stays yours.



Any Chat App

Talk to it on WhatsApp, Telegram, Discord, Slack, Signal, or iMessage. Works in DMs and group chats.



Persistent Memory

Remembers you and becomes uniquely yours. Your preferences, your context, your AI.



Browser Control

It can browse the web, fill forms, and extract data from any site.



Full System Access

Read and write files, run shell commands, execute scripts. Full access or sandboxed—your choice.



Skills & Plugins

Extend with community skills or build your own. It can even write its own.

Model Agnostic – works with Claude, ChatGPT, Gemini, Qwen, etc.

OPENCLAW USE-CASES



Productivity

- Jarvis like setup – No Computer? No problem
- Recurring tasks – Meeting setup , notes, scheduling
- Briefings – Weather, Emails summary
- Travel – Flight check-ins, stay etc

System Control

- Remote Device Control
- Browser Automation
- Cron Jobs & Webhooks
- Back up files

Work

- Code Review
- Content workflows – social media posting, newsletters etc

Build Your Own

OpenClaw uses a bunch of skills to do tasks

Skills – Pre-programmed instructions that tell what an LLM should do

Example: Calendar meeting scheduler, Email search etc

EXAMPLE: DINNER RESERVATIONS (PT 1)

...



Hi, I'd like



**Instructions & Text-to-Speech*

EXAMPLE: DINNER RESERVATIONS (PT 2)

...



Hi, I'd like



**Instructions & Text-to-Speech*

TOKENS AND APIS




What's a Token?

Every prompt you send and every response you get is measured in tokens.
1,000 tokens $\approx$ 750 words (roughly a full page of text).

<https://tiktokenizer.chatgptcn.com>

What's an API?

How software talk to each other. LLMs are invoked using APIs in a variety of applications. Claude leads the way in API usage



SECURITY



Threat Vectors

Prompt Injection

Malicious instructions hidden in emails & web pages trick the agent

Overprivileged Skills

Third-party plugins can silently exfiltrate data

Exposed Instances

Publicly accessible Gateway without proper auth

Autonomous Drift


Agent takes actions you never explicitly authorized

Best Practices

- Audit every skill before installing
- Use Tailscale / SSH tunnels with token auth
- Enable VirusTotal skill scanning
- Set guardrails: auto vs. confirm actions. Bind gateway to loopback
- Run on isolated infra (dedicated VPS)
- Review agent logs regularly

Real incidents: China banned OpenClaw in gov agencies | Cisco found skill exfiltration | Agent created dating profiles autonomously | Rogue agent published a hit piece

CORE ARCHITECTURE



DEPLOYMENT ARCHITECTURE / OPTIONS



Select the method that works for you and your setup



Local / Laptop

For development and testing only. Unreliable for production due to sleep, network changes, security, no static IP, etc.



macOS Virtual Machine

Sanboxed Virtual Machine (i.e. Isolated environment) via Lume on Apple Silicon.



Local / Dedicated Hardware

Mac mini or always-on Linux machine. You have full control, residential IP, and no cloud costs.



Virtual Private Server Cloud Virtual Machine

Always-on Linux server. This should work for most users who don't want to manage dedicated hardware.



Hybrid

Gateway on a Virtual Private Server, Browser/UI automation on a local machine. Always on reliability of VPS + Residential IP.



Note on Windows: OpenClaw is available on Native Windows, but official docs say that WSL2 is more stable.

<https://docs.openclaw.ai/platforms/windows>

KEY CONCEPTS & TERMS



API KEY / AUTH TOKENS

Function as passwords to the model provider (e.g., Anthropic, OpenAI, Google, etc.). When OpenClaw requests the model, it sends this key to prove it's authorized to use your account.

Treat API Keys like passwords.

SYSTEM PROMPT

Standing instruction that gets prepended to every single conversation before any user message arrives. It's where you define who the agent is, how it should behave, what it's allowed to do, and what tone it takes.

TEMPERATURE

A number between 0 and 1 that controls how predictable or creative the model's responses are. At 0, the model always picks the most statistically likely next token, fully deterministic and consistent.

Agents => 0 to 0.3 for reliability and consistency

STREAMING

Streaming means it sends tokens as they're generated, word by word, so the UI can display the response in real time as it's being written. Versus, a model generating the full response and sends it back in one block.

API KEYS / OAUTH TOKENS



To connect OpenClaw to a model, API Keys and OAuth Tokens act as passwords

API keys


Pay per token, separate from any subscription you already have. Recommended by OpenClaw docs for production.

Available across all model providers.


OAuth (Subscription Auth)

OpenClaw authenticates using your existing subscription and draws from its usage quota.

Availability depends on the provider (e.g. Claude Max)



In Terminal:
claude setup-token



INFRASTRUCTURE COSTS



Option	Monthly Cost	One-Time Cost	Notes
Local Laptop	\$0	—	Dev/testing only (not always-on, unreliable for production)
Local Mac mini / Dedicated	\$0/mo	~\$599+	Always-on, residential IP, best for browser automation (M4 M
Oracle Cloud ARM Free Tier	\$0	—	4 OCPU + 24GB RAM — best free cloud option; upgrade to PA avoid idle reclamation
Hetzner CAX11	~\$4–5/mo	—	Best value stable VPS; rising to ~\$5 from April 2026
DigitalOcean Droplet	~\$6/mo	—	Good baseline option
GCP e2-small	~\$12/mo	—	Minimum recommended for Docker builds
GCP e2-medium	~\$25/mo	—	Recommended for stable multi-channel deployments
Azure B2s	~\$30/mo	—	Most expensive mainstream option
MacStadium Hosted Mac	\$99+/mo	—	Only needed for cloud iMessage via BlueBubbles



WHAT'S NEXT (Depends on interest)



Claude Deep-dive

We cover the current frontier model landscape, then go deep on the skills that separate casual users from power users: structuring prompts like specifications, building persistent Projects, working with the API and connectors, and using Claude's platform features to create reusable workflows.

OpenClaw Setup / Hands-On

We get hands-on with a real OpenClaw deployment: installing the Gateway, connecting Telegram as a live channel, and sending your first message through your own self-hosted AI agent. By the end of the session, you'll have a running instance you own and control, ready to extend.

RELEVANT RESOURCES



- OpenClaw Technical Docs: <https://docs.openclaw.ai/>
- OpenClaw Install Guide: <https://docs.openclaw.ai/install>
- OpenClaw GitHub: <https://github.com/openclaw/openclaw>
- OpenClaw Awesome Skills: <https://github.com/VoltAgent/awesome-openclaw-skills>
- OpenClaw Awesome Use-cases: <https://github.com/hesamsheikh/awesome-openclaw-usecases>

Real OpenClaw  **Use Cases**



Format: We will open the OpenClaw docs together and walk through:

1. MAC Users: macOS VIRTUAL MACHINE (VM) LUME SETUP

1. Up to you: Google Cloud Project (GCP) + Virtual Private Server (VPS) Setup

Just in case you need help:

Please copy-paste the URL of the relevant technical doc into your LLM of choice to act as a troubleshooting partner for you.

MAC OS VIRTUAL MACHINE (VM) LUME SETUP

[https://docs.openclaw.ai/install
/macos-vm](https://docs.openclaw.ai/install/mac-os-vm)

GOOGLE CLOUD PROJECT (GCP) + VIRTUAL PRIVATE SERVER (VPS) SETUP

<https://docs.openclaw.ai/install/gcp>